


# Supervised learning

## Perceptron

## Section 1

### The perceptron (regression)

# Where we see a first neuron



$$h_w(x) = w_0 + w_1 \times x_1 + w_2 \times x_2 + \cdots + w_n \times x_n$$

# Characterisation

The elements of the  $\mathbf{w}$  vector are called the weights of the perceptron.


Given an input vector  $x$ , the output of the perceptron is a linear combination of its input.

$$h_w(x) = w_0 + w_1 \times x_1 + w_2 \times x_2 + \cdots + w_n \times x_n$$

The set of functions representable by a perceptron is the set of linear functions. That's our hypothesis space  $\mathcal{H}_{perc}$ .

# An even simpler dataset

X	Y
$m^2$	Price (€)
24	102 000
46	140 000
50	353 600
211	892 000
74	198 000



## Perceptron for predicting the price

We have a single feature ( $x_1$ : squared meters), thus a function representable by a perceptron would have the form:

$$h_w(sqm) = w_0 + w_1 \times x_1$$

where we can interpret:


- $w_0$  as the base price
- $w_1$  as the price per squared meters

# Perceptron for predicting the price

Making an educated guess we could set:

$$w_0 = 10000 \text{ (€)}$$

$$w_1 = 3000 \text{ (€/m}^2\text{)}$$



# Perceptron: prediction error


For an example  $(x, y)$  and predictor  $h_w$

prediction error =  $|y - h_w(x)|$  (length of red segments)

$L_2$  loss:  $L_2(y, \hat{y}) = (y - h_w(x))^2$  (squared error)

This leads to the empirical loss (for  $L_2$ ) over the entire dataset  $E$ :

$$EmpLoss_{L_2, E}(h_w) = \sum_{(x, y) \in E} \frac{(y - h_w(x))^2}{|E|}$$





## Finding the best hypothesis

I want the hypothesis  $\hat{h}^*$ , with the minimum empirical loss:

$$\hat{h}^* = \arg \min_{h_w \in \mathcal{H}_{perc}} EmpLoss_{L_2, E}(h_w)$$

For the perceptron, this means finding the best weights  $\hat{w}^*$  in the weight space.

$$\hat{w}^* = \arg \min_w EmpLoss_{L_2, E}(h_w)$$

Posing  $Loss(w) = EmpLoss_{L_2, E}(h_w)$ , we obtain:

$$\hat{w}^* = \arg \min_w Loss(w)$$

## Gradient: direction of steepest ascent

In any point  $w$  of the function, the gradient defines the direction of steepest ascent:

$$\vec{\nabla} g(w)$$

It can be computed from the partial derivatives:

$$\vec{\nabla} g(w) = \begin{bmatrix} \frac{\delta}{\delta w_0} g(w) \\ \frac{\delta}{\delta w_1} g(w) \\ \vdots \\ \frac{\delta}{\delta w_m} g(w) \end{bmatrix}$$




Figure: Gradient of  $f(x, y) = -(\cos^2(x) + \cos^2(y))^2$

## Gradient descent

From a point  $w$ , compute a new candidate  $w'$  by following the direction of steepest descent (opposite of the gradient).

$$w' = w - \alpha \times \vec{\nabla} g(w)$$

The distance traveled is parameterized by the **step size**  $\alpha$ .

Since the function is decreasing in this direction, there is a good chance that

$$g(w') < g(w)$$

# Gradient descent

Applying this repeatedly, we get the gradient descent algorithm:<sup>1</sup>

$w \leftarrow$  any value in the parameter space

**while** not converged **do**

$w \leftarrow w - \alpha \times \vec{\nabla} Loss(w)$

Typical convergence criteria: stop when the update did not provide an improvement for the last  $k$  iterations (e.g.  $k = 5$ ).

---

<sup>1</sup>Recal that  $Loss(w)$  is shortcut for  $EmpLoss_{L,E}(h_w)$

## Computing the gradient ( $L_2$ loss, single example)

Partial derivative of the  $L_2$  loss for a single example  $(x, y)$ :

$$\begin{aligned}\frac{\delta}{\delta w_i} \text{Loss}(w) &= \frac{\delta}{\delta w_i} (y - h_w(x))^2 \\ &= 2(y - h_w(x)) \times \frac{\delta}{\delta w_i} (y - h_w(x))\end{aligned}$$

Applied to our system with a single feature  $(x_1)$  we obtain:

$$\begin{aligned}\frac{\delta}{\delta w_0} \text{Loss}(w) &= -2(y - h_w(x)) \\ \frac{\delta}{\delta w_1} \text{Loss}(w) &= -2(y - h_w(x)) \times x_1\end{aligned}$$

## Update rules

Updating the weights based on a single example:<sup>2</sup>

$$\begin{aligned}w_0 &\leftarrow w_0 + \alpha \times (y - h_w(x)) \\w_1 &\leftarrow w_1 + \alpha \times (y - h_w(x)) \times x_1\end{aligned}$$


Updating the weights based on the entire training set  $E$ :

$$\begin{aligned}w_0 &\leftarrow w_0 + \alpha \times \sum_{(x,y) \in E} (y - h_w(x)) \\w_1 &\leftarrow w_1 + \alpha \times \sum_{(x,y) \in E} (y - h_w(x)) \times x_1\end{aligned}$$

---


<sup>2</sup>Note that the  $-2$  factor from the previous equation is included in  $\alpha$  term.

# Updating our initial guess<sup>3</sup>



$$w = [10000, 3000]$$

<sup>3</sup>With  $\alpha = 10^{-5}$



$$w' = [10010.53, 4343.82]$$


# Result of the gradient descent

Repeating the gradient descent step, we eventually converge to our best solution.

$$\hat{w}^* = [9947.29, 4235.02]$$

I should sell my apartment for:

$$9947.29 + 4235.02 \times 165 = 708725\text{€}$$





# Gradient descent (applied to a perceptron with $L_2$ loss)

$w \leftarrow$  any value in the parameter space

**while** not converged **do**

$$w_0 \leftarrow w_0 + \alpha \times \sum_{(x,y) \in E} (y - h_w(x))$$

**for**  $i \in 1 \dots n$  **do**

$$w_i \leftarrow w_i + \alpha \times \sum_{(x,y) \in E} (y - h_w(x)) \times x_i$$

## Representation trick

In the previous slides, we always had to deal with the  $w_0$  weight specially because it has no corresponding feature.

We can define an artificial feature  $x_0$  that always has the value 1. And reformulate  $h_w$ :

$$h_w(x) = \sum_{i \in [0, m]} w_i \times x_i$$

$$h_w(x) = w \cdot x \quad (\text{dot product})$$

and the update rule (for  $L_2$  loss):

$$w_i \leftarrow w_i + \alpha \times \sum_{(x, y) \in E} (y - h_w(x)) \times x_i$$

## Section 2

### A perceptron for classification

## A classification problem

I now want to buy a new apartment to replace the one I just sold. To be reactive I built an automated system that sends me any new announce of an apartment for sale.


- Problem: there are dozens of announces every day and I don't have time to look at them all.
- Solution: build an AI system that will predict whether I will be interested in a particular apartment based on a few of its features. If it predicts that I am not interested, it will discard the announce.

## A classification problem: dataset


So far, I collect the following information stating whether an announce that I previously saw was interesting.

X				Y
$m^2$	Num Rooms	Floor	Price (€)	Interesting
24	1	4	102 000	true
46	3	2	140 000	false
50	3	6	353 600	false
211	5	3	892 000	true
74	3	1	198 000	true

# A classification problem: dataset visualization




## Our previous perceptron



$$h_w(x) = w_0 + w_1 \times x_1 + w_2 \times x_2 + \cdots + w_n \times x_n$$

$$h_w(x) = w \cdot x$$

# Perceptron for classification




$$h_w(x) = \text{Step}(w \cdot x)$$

where  $\text{Step}(z) = 1$  if  $z \geq 0$  and 0 otherwise




# The perceptron as a linear classifier

The perceptron defines a **decision boundary** that separates two classes.



Linearly separable

Perfectly classifiable by a perceptron



**Not** linearly separable

**Not** perfectly classifiable by a perceptron

# The step function

$$h_w(x) = \text{Step}(w \cdot x)$$


where  $\text{Step}(z) = 1$  if  $z \geq 0$  and 0 otherwise

We now have a function that we could train in order to output:

- 1 if the example is in the class (interesting)
- 0 otherwise (not interesting)

Problem: the function is:

- non-differentiable in 0
- the gradient is 0 everywhere else



$\text{Step}(z)$

# The perceptron learning rule

Nevertheless, an rule was proposed the **perceptron update rule** (here for a single example  $(x, y)$ ):



$$w_i \leftarrow w_i + \alpha \times (y - h_w(x)) \times x_i$$

which is identical to the update rule for linear regression (for  $L_2$ ).

The rule is show to converge to a solution when the data is linearly separable.

## The perceptron learning rule (under non separable data)

However the perceptron learning rule is unstable when the data is not linearly separable:



## Replacing the step function


Turns out we can replace the step function with one with nicer properties.

$$\textit{Logistic}(z) = \frac{1}{1 + e^{-z}}$$

and redefine our hypothesis function:

$$h_w(x) = \textit{Logistic}(w \cdot x) = \frac{1}{1 + e^{-w \cdot x}}$$

Often called the **logistic regression**.



## Back on track: gradient descent

*Logistic*( $x$ ) is differentiable on  $] -\infty, \infty[$ . This allows us to reuse gradient descent for training:



$$w \leftarrow w - \alpha \times \vec{\nabla} Loss(w)$$

For an  $L_2$  loss we obtain the update rule:

$$w_i \leftarrow w_i + \alpha(y - h_w(x)) \times h_w(x) \times (1 - h_w(x)) \times x_i$$

## Training the logistic regression (under non-separable data)

Compared to the step function, the logistic regression tends to converge more quickly and reliably in the presence of noisy and non-separable data.



## Section 3

### Synthesis




# Synthesis

We saw two classes of perceptrons:

- linear regressor
- linear classifier

Both can be trained with **gradient descent** in attempt to **minimize the loss**.

In the next course, the perceptron will be a **neural unit** in a **neural network**.



# Exercises

- For each of the following datasets, propose a perceptron that would correctly classify it.

X	Y
-1	Flower
2	Cat
1	Flower
0	Flower

X	Y
-1	Flower
2	Cat
1	Cat
0	Cat

X	Y
-1	Flower
2	Flower
1	Cat
0	Cat

## Exercises

- 1 You have  $N$  examples in your dataset, each with  $M$  features. Give an estimate of the computational cost of a single update step. Does it scale to large-scale datasets (e.g.  $N = 10^5, M = 10^4$ )
- 2 For the linear regression (regression perceptron), are we guaranteed to find the optimal weights with gradient descent?
- 3 What's a reasonable loss for spam detection? (you can make some assumptions about the proportion of spam)
- 4 What's the update formula of a regression perceptron using the  $L_1$  loss?